

```

1  store_prod:
2      movq    %rdx, %rax
3      cqto
4      movq    %rsi, %rcx
5      sarq    $63, %rcx
6      imulq   %rax, %rcx
7      imulq   %rsi, %rdx
8      addq    %rdx, %rcx
9      mulq    %rsi
10     addq    %rcx, %rdx
11     movq    %rax, (%rdi)
12     movq    %rdx, 8(%rdi)
13     ret

```

This code uses three multiplications for the multiprecision arithmetic required to implement 128-bit arithmetic on a 64-bit machine. Describe the algorithm used to compute the product, and annotate the assembly code to show how it realizes your algorithm. *Hint:* When extending arguments of x and y to 128 bits, they can be rewritten as $x = 2^{64} \cdot x_h + x_l$ and $y = 2^{64} \cdot y_h + y_l$, where x_h, x_l, y_h , and y_l are 64-bit values. Similarly, the 128-bit product can be written as $p = 2^{64} \cdot p_h + p_l$, where p_h and p_l are 64-bit values. Show how the code computes the values of p_h and p_l in terms of x_h, x_l, y_h , and y_l .

3.60 ♦♦

Consider the following assembly code:

```

long loop(long x, int n)
x in %rdi, n in %esi
1  loop:
2      movl    %esi, %ecx
3      movl    $1, %edx
4      movl    $0, %eax
5      jmp     .L2
6  .L3:
7      movq    %rdi, %r8
8      andq    %rdx, %r8
9      orq     %r8, %rax
10     salq    %cl, %rdx
11  .L2:
12     testq   %rdx, %rdx
13     jne     .L3
14     rep; ret

```

The preceding code was generated by compiling C code that had the following overall form:

```

1  long loop(long x, long n)
2  {
3      long result = _____;
4      long mask;
5      for (mask = _____; mask _____ ; mask = _____ ) {
6          result |= _____ ;
7      }
8      return result;
9  }

```

Your task is to fill in the missing parts of the C code to get a program equivalent to the generated assembly code. Recall that the result of the function is returned in register %rax. You will find it helpful to examine the assembly code before, during, and after the loop to form a consistent mapping between the registers and the program variables.

- A. Which registers hold program values `x`, `n`, `result`, and `mask`?
- B. What are the initial values of `result` and `mask`?
- C. What is the test condition for `mask`?
- D. How does `mask` get updated?
- E. How does `result` get updated?
- F. Fill in all the missing parts of the C code.

3.61 ♦♦

In Section 3.6.6, we examined the following code as a candidate for the use of conditional data transfer:

```

long cread(long *xp) {
    return (xp ? *xp : 0);
}

```

We showed a trial implementation using a conditional move instruction but argued that it was not valid, since it could attempt to read from a null address.

Write a C function `cread_alt` that has the same behavior as `cread`, except that it can be compiled to use conditional data transfer. When compiled, the generated code should use a conditional move instruction rather than one of the jump instructions.

3.62 ♦♦

The code that follows shows an example of branching on an enumerated type value in a switch statement. Recall that enumerated types in C are simply a way to introduce a set of names having associated integer values. By default, the values assigned to the names count from zero upward. In our code, the actions associated with the different case labels have been omitted.

```

1  /* Enumerated type creates set of constants numbered 0 and upward */
2  typedef enum {MODE_A, MODE_B, MODE_C, MODE_D, MODE_E} mode_t;
3
4  long switch3(long *p1, long *p2, mode_t action)
5  {
6      long result = 0;
7      switch(action) {
8          case MODE_A:
9
10         case MODE_B:
11
12         case MODE_C:
13
14         case MODE_D:
15
16         case MODE_E:
17
18         default:
19
20     }
21     return result;
22 }

```

The part of the generated assembly code implementing the different actions is shown in Figure 3.52. The annotations indicate the argument locations, the register values, and the case labels for the different jump destinations.

Fill in the missing parts of the C code. It contained one case that fell through to another—try to reconstruct this.

3.63 ♦♦

This problem will give you a chance to reverse engineer a `switch` statement from disassembled machine code. In the following procedure, the body of the `switch` statement has been omitted:

```

1  long switch_prob(long x, long n) {
2      long result = x;
3      switch(n) {
4          /* Fill in code here */
5
6      }
7      return result;
8  }

```

```

    p1 in %rdi, p2 in %rsi, action in %edx
1   .L8:                                MODE_E
2       movl    $27, %eax
3       ret
4   .L3:                                MODE_A
5       movq    (%rsi), %rax
6       movq    (%rdi), %rdx
7       movq    %rdx, (%rsi)
8       ret
9   .L5:                                MODE_B
10      movq    (%rdi), %rax
11      addq    (%rsi), %rax
12      movq    %rax, (%rdi)
13      ret
14   .L6:                                MODE_C
15      movq    $59, (%rdi)
16      movq    (%rsi), %rax
17      ret
18   .L7:                                MODE_D
19      movq    (%rsi), %rax
20      movq    %rax, (%rdi)
21      movl    $27, %eax
22      ret
23   .L9:                                default
24      movl    $12, %eax
25      ret

```

Figure 3.52 Assembly code for Problem 3.62. This code implements the different branches of a switch statement.

Figure 3.53 shows the disassembled machine code for the procedure.

The jump table resides in a different area of memory. We can see from the indirect jump on line 5 that the jump table begins at address 0x4006f8. Using the GDB debugger, we can examine the six 8-byte words of memory comprising the jump table with the command `x/6gx 0x4006f8`. GDB prints the following:

```

(gdb) x/6gx 0x4006f8
0x4006f8:  0x00000000004005a1  0x00000000004005c3
0x400708:  0x00000000004005a1  0x00000000004005aa
0x400718:  0x00000000004005b2  0x00000000004005bf

```

Fill in the body of the switch statement with C code that will have the same behavior as the machine code.

```

long switch_prob(long x, long n)
x in %rdi, n in %rsi
1 0000000000400590 <switch_prob>:
2 400590: 48 83 ee 3c          sub    $0x3c,%rsi
3 400594: 48 83 fe 05          cmp    $0x5,%rsi
4 400598: 77 29                ja     4005c3 <switch_prob+0x33>
5 40059a: ff 24 f5 f8 06 40 00 jmpq   *0x4006f8(,%rsi,8)
6 4005a1: 48 8d 04 fd 00 00 00 lea     0x0(,%rdi,8),%rax
7 4005a8: 00
8 4005a9: c3                  retq
9 4005aa: 48 89 f8            mov    %rdi,%rax
10 4005ad: 48 c1 f8 03         sar    $0x3,%rax
11 4005b1: c3                  retq
12 4005b2: 48 89 f8            mov    %rdi,%rax
13 4005b5: 48 c1 e0 04         shl    $0x4,%rax
14 4005b9: 48 29 f8            sub    %rdi,%rax
15 4005bc: 48 89 c7            mov    %rax,%rdi
16 4005bf: 48 0f af ff         imul   %rdi,%rdi
17 4005c3: 48 8d 47 4b         lea     0x4b(%rdi),%rax
18 4005c7: c3                  retq

```

Figure 3.53 Disassembled code for Problem 3.63.

3.64 ♦♦♦

Consider the following source code, where R , S , and T are constants declared with `#define`:

```

1 long A[R][S][T];
2
3 long store_ele(long i, long j, long k, long *dest)
4 {
5     *dest = A[i][j][k];
6     return sizeof(A);
7 }

```

In compiling this program, gcc generates the following assembly code:

```

long store_ele(long i, long j, long k, long *dest)
i in %rdi, j in %rsi, k in %rdx, dest in %rcx
1 store_ele:
2 leaq    (%rsi,%rsi,2), %rax
3 leaq    (%rsi,%rax,4), %rax
4 movq    %rdi, %rsi
5 salq    $6, %rsi
6 addq    %rsi, %rdi
7 addq    %rax, %rdi

```