

- A. What is the total number of writes?
- B. What is the total number of writes that hit in the cache?
- C. What is the hit rate?

**6.40 ♦**

Given the assumptions in Problem 6.38, determine the cache performance of the following code:

```

1      for (i = 15; i >= 0; i--) {
2          for (j = 15; j >= 0; j--) {
3              square[i][j].y = 1;
4          }
5      }
6      for (i = 15; i >= 0; i--) {
7          for (j = 15; j >= 0; j--) {
8              square[i][j].c = 0;
9              square[i][j].m = 0;
10             square[i][j].k = 0;
11         }
12     }
```

- A. What is the total number of writes?
- B. What is the total number of writes that hit in the cache?
- C. What is the hit rate?

**6.41 ♦♦**

You are writing a new 3D game that you hope will earn you fame and fortune. You are currently working on a function to blank the screen buffer before drawing the next frame. The screen you are working with is a  $640 \times 480$  array of pixels. The machine you are working on has a 32 KB direct-mapped cache with 8-byte lines. The C structures you are using are as follows:

```

1  struct pixel {
2      char r;
3      char g;
4      char b;
5      char a;
6  };
7
8  struct pixel buffer[480][640];
9  int i, j;
10 char *cptr;
11 int *iptr;
```

Assume the following:

- `sizeof(char) = 1` and `sizeof(int) = 4`.

- buffer begins at memory address 0.
- The cache is initially empty.
- The only memory accesses are to the entries of the array buffer. Variables *i*, *j*, *cptr*, and *iptr* are stored in registers.

What percentage of writes in the following code will hit in the cache?

```

1      for (j = 639; j >= 0; j--) {
2          for (i = 479; i >= 0; i--){
3              buffer[i][j].r = 0;
4              buffer[i][j].g = 0;
5              buffer[i][j].b = 0;
6              buffer[i][j].a = 0;
7          }
8      }
```

#### 6.42 ♦♦

Given the assumptions in Problem 6.41, what percentage of writes in the following code will hit in the cache?

```

1      char *cptr = (char *) buffer;
2      for (; cptr < (((char *) buffer) + 640 * 480 * 4); cptr++)
3          *cptr = 0;
```

#### 6.43 ♦♦

Given the assumptions in Problem 6.41, what percentage of writes in the following code will hit in the cache?

```

1      int *iptr = (int *)buffer;
2      for (; iptr < ((int *)buffer + 640*480); iptr++)
3          *iptr = 0;
```

#### 6.44 ♦♦♦

Download the mountain program from the CS:APP Web site and run it on your favorite PC/Linux system. Use the results to estimate the sizes of the caches on your system.

#### 6.45 ♦♦♦♦

In this assignment, you will apply the concepts you learned in Chapters 5 and 6 to the problem of optimizing code for a memory-intensive application. Consider a procedure to copy and transpose the elements of an  $N \times N$  matrix of type *int*. That is, for source matrix *S* and destination matrix *D*, we want to copy each element  $s_{i,j}$  to  $d_{j,i}$ . This code can be written with a simple loop,

```

1  void transpose(int *dst, int *src, int dim)
2  {
3      int i, j;
4
```

```

5      for (i = 0; i < dim; i++)
6          for (j = 0; j < dim; j++)
7              dst[j*dim + i] = src[i*dim + j];
8  }
```

where the arguments to the procedure are pointers to the destination (`dst`) and source (`src`) matrices, as well as the matrix size  $N$  (`dim`). Your job is to devise a transpose routine that runs as fast as possible.

#### 6.46 ♦♦♦♦

This assignment is an intriguing variation of Problem 6.45. Consider the problem of converting a directed graph  $g$  into its undirected counterpart  $g'$ . The graph  $g'$  has an edge from vertex  $u$  to vertex  $v$  if and only if there is an edge from  $u$  to  $v$  or from  $v$  to  $u$  in the original graph  $g$ . The graph  $g$  is represented by its *adjacency matrix*  $G$  as follows. If  $N$  is the number of vertices in  $g$ , then  $G$  is an  $N \times N$  matrix and its entries are all either 0 or 1. Suppose the vertices of  $g$  are named  $v_0, v_1, v_2, \dots, v_{N-1}$ . Then  $G[i][j]$  is 1 if there is an edge from  $v_i$  to  $v_j$  and is 0 otherwise. Observe that the elements on the diagonal of an adjacency matrix are always 1 and that the adjacency matrix of an undirected graph is symmetric. This code can be written with a simple loop:

```

1  void col_convert(int *G, int dim) {
2      int i, j;
3
4      for (i = 0; i < dim; i++)
5          for (j = 0; j < dim; j++)
6              G[j*dim + i] = G[j*dim + i] || G[i*dim + j];
7  }
```

Your job is to devise a conversion routine that runs as fast as possible. As before, you will need to apply concepts you learned in Chapters 5 and 6 to come up with a good solution.

## Solutions to Practice Problems

### Solution to Problem 6.1 (page 620)

The idea here is to minimize the number of address bits by minimizing the aspect ratio  $\max(r, c) / \min(r, c)$ . In other words, the squarer the array, the fewer the address bits.

Organization	$r$	$c$	$b_r$	$b_c$	$\max(b_r, b_c)$
$16 \times 1$	4	4	2	2	2
$16 \times 4$	4	4	2	2	2
$128 \times 8$	16	8	4	3	4
$512 \times 4$	32	16	5	4	5
$1,024 \times 4$	32	32	5	5	5